

Yelp Dataset Review

John Antony

March 4, 2017

Reading review data from JSON file

```
library(readr)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
reviews_file_path = "/Users/JohnAntony/Desktop/Main/Applications/R/MachineLearning/Yelp/yelp_dataset_challenge_roun9/yelp_academic_dataset_review.json"

review_lines <- read_lines(reviews_file_path, n_max = 100000, progress = FALSE)

library(stringr)
library(jsonlite)

reviews_combined <- str_c("[", str_c(review_lines, collapse = ", "), "]")

reviews <- fromJSON(reviews_combined) %>%
  flatten() %>%
  tbl_df()
#head(reviews)
remove(review_lines, reviews_combined)
```

Reading Business data from JSON File

```
business_file_path = "/Users/JohnAntony/Desktop/Main/Applications/R/MachineLearning/Yelp/yelp_dataset_challenge_roun9/yelp_academic_dataset_business.json"

business_lines <- read_lines(business_file_path, n_max = 100000, progress = FALSE)

business_combined <- str_c("[", str_c(business_lines, collapse = ", "), "]")

business <- fromJSON(business_combined) %>%
  flatten() %>%
  tbl_df()

#head(business)
business[order(business$business_id),]
```

```
## # A tibble: 100,000 × 16
##           business_id          name
##           <chr>             <chr>
## 1  __1uG7MLxWGFiv2fCGPiQQ  SpinalWorks Chiropractic
## 2  __8j8yhsmE98wNWHJNyAgw    Urawa Sushi
## 3  __blIPRrsfEoaiOSPj1olQ    Property Frameworks
## 4  __bqGGnOjtY9eEhrZAUsGA    Galangal Thai Fusion
## 5  __CQ2SE4NXFFjYfrB_TJ6w St. Gabriel Medical Clinic
## 6  __D6AVR_hLpW_bott0-upA    Skinapeel Beauty
## 7  __FFoyg0XmJluBBNE0QP0w    Better Health Solutions
## 8  __fMLrmv9M1_W4kBvR2VnQ    Dairy Queen
## 9  __G0Ug3CK2yCDdQLYpd0ww    LV spa
## 10 __H_61gpm7eViPMbWxPZSg    Subway
## # ... with 99,990 more rows, and 14 more variables: neighborhood <chr>,
## #   address <chr>, city <chr>, state <chr>, postal_code <chr>,
## #   latitude <dbl>, longitude <dbl>, stars <dbl>, review_count <int>,
## #   is_open <int>, attributes <list>, categories <list>, hours <list>,
## #   type <chr>
```

```
remove(business_combined,business_lines)
```

Cleanup

```
rest_reviews = aggregate(text ~ business_id, data = reviews, paste, collapse = ",")
rest_reviews$text = tolower(rest_reviews$text)
#head(rest_reviews)
```

Defining functions for Sentiment Scoring and Pulling positive and negative words

```
score.sentiment = function(sentences, pos.words, neg.words, .progress='none')
{
  require(plyr)
  require(stringr)

  # we got a vector of sentences. plyr will handle a list or a vector as an "l" for us
  # we want a simple array of scores back, so we use "l" + "a" + "ply" = laply:
  scores = laply(sentences, function(sentence, pos.words, neg.words) {

    # clean up sentences with R's regex-driven global substitute, gsub():
    sentence = gsub('[:punct:]', '', sentence)
    sentence = gsub('[:cntrl:]', '', sentence)
    sentence = gsub('\\d+', '', sentence)
    # and convert to lower case:
    sentence = tolower(sentence)

    # split into words. str_split is in the stringr package
    word.list = str_split(sentence, '\\s+')
    # sometimes a list() is one level of hierarchy too much
    words = unlist(word.list)

    # compare our words to the dictionaries of positive & negative terms
    pos.matches = match(words, pos.words)
    neg.matches = match(words, neg.words)

    # match() returns the position of the matched term or NA
    # we just want a TRUE/FALSE:
    pos.matches = !is.na(pos.matches)
    neg.matches = !is.na(neg.matches)

    # and conveniently enough, TRUE/FALSE will be treated as 1/0 by sum():
    score = sum(pos.matches) - sum(neg.matches)

    return(score)
  }, pos.words, neg.words, .progress=.progress )

  scores.df = data.frame(score=scores, text=sentences)
  return(scores.df)
}
```

```

HIDict = readLines("/Users/JohnAntony/Desktop/Main/Applications/R/MachineLearning/data_files/inqdict.txt")
dict_pos = HIDict[grepl("Pos",HIDict)]
poswords = NULL
for (s in dict_pos) {
  s = strsplit(s,"#")[[1]][1]
  poswords = c(poswords,strsplit(s," ")[[1]][1])
}
dict_neg = HIDict[grepl("Neg",HIDict)]
negwords = NULL
for (s in dict_neg) {
  s = strsplit(s,"#")[[1]][1]
  negwords = c(negwords,strsplit(s," ")[[1]][1])
}
poswords = tolower(poswords)
negwords = tolower(negwords)
pos.words = unique(poswords)
neg.words = unique(negwords)

```

Sentiment Score for Review Texts

```
score = score.sentiment(rest_reviews$text, pos.words, neg.words)
```

```
## Loading required package: plyr
```

```
## -----
```

```

## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)

```

```
## -----
```

```

##
## Attaching package: 'plyr'

```

```
## The following objects are masked from 'package:dplyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize
```

```
rest_reviews_score = cbind(rest_reviews$business_id, score$score)
colnames(rest_reviews_score) <- c("business_id", "SentimentScore")
rest_reviews_score = as.data.frame(rest_reviews_score)

#rest_reviews_score
```

Creating Dataframe for regression

```
suppressMessages(library(dplyr))
#business_score_rating = merge(x = rest_reviews_score[,c("SentimentScore")], y = business[,c("stars")], by.res
t_reviews_score='business_id', by.business='business_id', all.x = TRUE)

business_score_rating = dplyr::left_join(rest_reviews_score, business, by = "business_id")
```

```
## Warning in left_join_impl(x, y, by$x, by$y, suffix$x, suffix$y): joining
## character vector and factor, coercing into character vector
```

```
business_score_rating = dplyr::select(business_score_rating, SentimentScore, stars)
business_score_rating$SentimentScore <- as.numeric(as.character(business_score_rating$SentimentScore))
#head(business_score_rating)
```

Regression of Sentiment Score with Business Score

```
res = lm(business_score_rating$SentimentScore ~ business_score_rating$stars)
print(summary(res))
```

```
##
## Call:
## lm(formula = business_score_rating$SentimentScore ~ business_score_rating$stars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -109.29  -21.40  -10.18    8.82   605.77
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      11.6247     1.4733    7.89 3.32e-15 ***
## business_score_rating$stars  5.8879     0.3938   14.95 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 40.28 on 10275 degrees of freedom
## (4399 observations deleted due to missingness)
## Multiple R-squared:  0.02129,    Adjusted R-squared:  0.02119
## F-statistic: 223.5 on 1 and 10275 DF,  p-value: < 2.2e-16
```

Creating a Network Cloud on Positive and Negative Sentiments

```
reviews_rest <- subset(reviews, business_id=="GdCIMZ9BTT4ywETWcByfJA")
#tail(names(sort(table(reviews_rest$business_id))), 1)
#head(reviews_rest)
```

Getting the bigrams

```
library(dplyr)
library(tidytext)
library(tidyr)

reviews_text = subset(reviews_rest, select = c(text) )

reviews_bigrams = reviews_text %>%
  unnest_tokens(bigram, text, token = "ngrams", n = 2)
#reviews_bigrams
```

Using dplyr Count for counting the occurrence of bigrams

```
#reviews_bigrams %>%
#   dplyr::count(bigram, sort = TRUE)
```

Bigram Cleaning for Text analysis

```
bigrams_separated <- reviews_bigrams %>%
  separate(bigram, c("word1", "word2"), sep = " ")

bigrams_filtered <- bigrams_separated %>%
  filter(!word1 %in% stop_words$word) %>%
  filter(!word2 %in% stop_words$word) %>%
  filter(word1 %in% neg.words)

bigram_counts <- bigrams_filtered %>%
  dplyr::count(word1, word2, sort = TRUE)
remove(reviews_bigrams)
```

Using igraph to discover the network graph

```
library(igraph)
```



```
##  
## Attaching package: 'igraph'
```

```
## The following objects are masked from 'package:tidyr':  
##  
## %>%, crossing
```

```
## The following object is masked from 'package:stringr':  
##  
## %>%
```

```
## The following objects are masked from 'package:dplyr':  
##  
## %>%, as_data_frame, groups, union
```

```
## The following objects are masked from 'package:stats':  
##  
## decompose, spectrum
```

```
## The following object is masked from 'package:base':  
##  
## union
```

```
bigram_graph <- bigram_counts %>%  
  graph_from_data_frame()  
  
bigram_graph
```

```
## IGRAPH DN-- 103 71 --
## + attr: name (v/c), n (e/n)
## + edges (vertex names):
## [1] club      ->card      fire      ->roasted  service  ->industry
## [4] awful     ->lot       awful     ->simply   awkward  ->mishmash
## [7] bad       ->experience bad       ->location bad       ->mall
## [10] bad      ->reminded bad      ->reviews bit       ->crazy
## [13] bit      ->messy     bit      ->pricey  bit       ->soft
## [16] black    ->cardinals black    ->hummers black     ->yukons
## [19] bland    ->thai      club     ->chefs   cold      ->cut
## [22] cut      ->style     damn     ->chocolate difficult->time
## + ... omitted several edges
```

```
library(ggraph)
```

```
## Loading required package: ggplot2
```

```
set.seed(2017)
```

```
ggraph(bigram_graph, layout = "fr") +
  geom_edge_link() +
  geom_node_point() +
  geom_node_text(aes(label = name), vjust = 1, hjust = 1)
```

